Response Propensity Estimation for Combining Probability and Non-Probability Samples with LBBNN

Aliaksandr Hubin & Lars Skaaret-Lund (in collaboration with Vladislav Beresovsky, Julie Gershunskaya

2026-11-13

Contents

1 Executive Summary **2**

2 Theoretical Foundations **2**

2.1 The Challenge of Non-Probability Convenience Data 2

2.2 The Five Estimation Formulations (Beresovsky, Gershunskaya & Savitsky, 2026) 3

2.3 Self-Normalizing Hájek Estimator 7

2.4 Variational ELBO Calibration & Protection Against Cold Posteriors 7

3 Package Setup & Environment Verification **7**

4 Simulation Study 0: Exact Linear Baseline (Beresovsky et al., 2026) **8**

4.1 High Common Support Benchmark 11

4.2 Low Common Support Benchmark 12

5 Simulation Study 1: High Common Support with Linear Propensity & Noise Covariates **13**

5.1 Fitting Propensity Models (LBBNN vs All 5 GLM Baselines) 14

5.2 Sparse Variable Selection & Active Graph Pruning 15

5.3 Finite Population Mean Recovery 16

6 Simulation Study 2: High Common Support with Non-Linear Propensity & Interaction **17**

6.1 Fitting Propensity Models under High Support Non-Linearity 18

6.2 Variable Selection under Non-Linear Interaction 19

6.3 Finite Population Mean Recovery 19

7 Simulation Study 3: Low Common Support with Non-Linear Propensity & Interaction **21**

7.1 Model Fitting under Low Common Support 22

7.2 Performance Comparison: Kish's ESS & The Breakdown of Pseudo-Likelihood 22

7.3 Propensity Recovery Diagnostic: 2x2 Comparison Grid 24

7.4 Empirical Demonstration: Vectorized Linkage Probabilities & Doubly Robust Track 2 -> Track 1 Bridge 27

8 Synthesis & Comparative Discussion Across Simulation Regimes **31**

8.1 When the True Propensity is Linear (Studies 0 & 1) 31

8.2 When the Propensity is Non-Linear with Interactions under High Common Support (Study 2) 31

8.3 When Non-Linearity Meets Low Common Support (Study 3) 32

9 Collaborator API Reference **32**

10 References **32**

11 Session Information **33**

1 Executive Summary

In federal and official statistics, the combination of a large **non-probability ("convenience") sample** (such as commercial web panels or administrative filings) with a smaller, high-quality **probability reference survey** offers substantial promise for reducing survey costs and respondent burden while expanding data coverage. However, unweighted convenience data suffer from severe **selection bias** (Meng, 2018).

This report implements and benchmarks the **Latent Binary Bayesian Neural Network (`LBBNN`)** framework for estimating unknown convenience participation propensities $\pi_{ci} = \mathbb{P}(I_{ci} = 1 \mid \mathbf{x}_i)$ under five composite likelihood and pseudo-likelihood formulations:

1. `STACK`: Composite Bernoulli likelihood on the unlinked stacked sample $S = S_c \oplus S_r$ ($\pi_{zi} = \frac{\pi_{ci}}{\pi_{ci}+\pi_{ri}}$).
2. `LINK1`: Deterministically linked union sample $S = S_c \cup (S_r \setminus S_c)$ ($\pi_{zi} = \frac{\pi_{ci}}{\pi_{ci}+\pi_{ri}-\pi_{ci}\pi_{ri}}$).
3. `LINK2`: Deterministically linked stacked sample $S = S_c \oplus S_r$ ($\pi_{zi} = \frac{\pi_{ci}(1+\pi_{ri})}{\pi_{ci}+\pi_{ri}}$).
4. `MIP` **(Missing Information Principle)**: Probabilistically linked stacked sample with linkage informativeness parameter $p_{\text{link}} \in [0, 1]$.
5. `PSEUDO`: Pseudo-log-likelihood approach of Chen, Li & Wu (2020) utilizing survey design weights $w_{ri} = 1/\pi_{ri}$ on S_r.

Through automated **variational spike-and-slab edge selection**, `LBBNN` simultaneously: - Captures complex **non-linear propensity curves and multi-way interactions** without manual polynomial expansions. - Prunes **uninformative noise covariates** from the network architecture. - Prevents the catastrophic **variance and weight explosion** that afflicts classical pseudo-likelihood under low common support.

---

2 Theoretical Foundations

2.1 The Challenge of Non-Probability Convenience Data

Let $U = \{1, \ldots, N\}$ denote a finite target population. Each unit $i \in U$ is associated with: - A vector of auxiliary covariates $\mathbf{x}_i \in \mathbb{R}^p$. - An outcome of interest $y_i \in \mathbb{R}$. - An unobserved convenience sample participation indicator $I_{ci} \in \{0, 1\}$, with unknown participation propensity $\pi_{ci} = \mathbb{P}(I_{ci} = 1 \mid \mathbf{x}_i) > 0$. - A probability survey inclusion indicator $I_{ri} \in \{0, 1\}$, with known design selection probability $\pi_{ri} = \mathbb{P}(I_{ri} = 1 \mid \mathbf{x}_i) > 0$.

If units with particular values of $\mathbf{x}_i$ or y_i systematically self-select into the convenience dataset, the naive sample mean is severely biased:

$$\text{Bias}(\bar{y}_c) = \frac{\text{Cov}_U(I_c, Y)}{\bar{\pi}_c}$$

As Meng (2018) proved, this selection bias scales with $\sqrt{N}$ rather than sample size n_c: even a massive convenience dataset comprising 80% or 90% of the population can produce wildly misleading population estimates if $\text{Corr}(I_c, Y) \neq 0$.

2.2 The Five Estimation Formulations (Beresovsky, Gershunskaya & Savitsky, 2026)

2.2.1 The Fundamental Identification Problem & The Conditioning Trick

In standard supervised machine learning, fitting a binary classification model $\mathbb{P}(I_{ci} = 1 \mid \mathbf{x}_i)$ requires a training sample of both positive ($I_{ci} = 1$) and negative ($I_{ci} = 0$) units drawn from the population. In the non-probability survey setting, however, we encounter a **fundamental identification barrier**:

1. We observe auxiliary covariates $\mathbf{x}_i$ only for units that participated in the convenience sample $S_c = \{i \in U : I_{ci} = 1\}$ and units selected into the reference probability survey $S_r = \{i \in U : I_{ri} = 1\}$.
2. For all non-sampled population units ($i \notin S_c \cup S_r$), the covariates $\mathbf{x}_i$ are entirely unobserved.

Consequently, one cannot directly fit a standard binary regression for I_{ci} over U. Beresovsky, Gershunskaya & Savitsky (2026) resolve this challenge through a **conditioning trick**: by combining S_c and S_r into a combined sample S and conditioning on the event that a unit was observed in S, the known reference design probability π_{ri} acts as a **structural offset** that anchors the neural network's unconstrained logits $\eta_\theta(\mathbf{x}_i)$ to the true population scale, allowing exact recovery of $\pi_{ci}(\mathbf{x}_i; \theta) = \sigma(\eta_\theta(\mathbf{x}_i))$.

Table 1: Overview of the Five Sample Combination Mechanisms (where $D_i = \mathbb{I}(i \in S_c)$ denotes sample origin)

Estimator	Sample Domain	Target Indicator & Linkage	Composite Probability / Mechanism
STACK	$S_c \oplus S_r$	Unlinked: $I_{zi} = D_i$	$\pi_{zi} = \frac{\pi_{ci}}{\pi_{ci}+\pi_{ri}}$
LINK1	$S_c \cup (S_r \setminus S_c)$	Dedup. Union: $I_{zi} = I_{ci}$ ($S_r \cap S_c$ dropped)	$\pi_{zi} = \frac{\pi_{ci}}{\pi_{ci}+\pi_{ri}-\pi_{ci}\pi_{ri}}$
LINK2	$S_c \oplus S_r$	Linked Stacked: $I_{zi} = D_i + (1 - D_i)I_{ci}$	$\pi_{zi} = \frac{\pi_{ci}(1+\pi_{ri})}{\pi_{ci}+\pi_{ri}}$
MIP	$S_c \oplus S_r$	Prob. Linked ($p_{\text{link}} \in [0, 1]$): $\tilde{I}_{zi} = D_i + (1 - D_i)\tilde{p}_i$	$\pi_{zi} = \frac{\pi_{ci}(1+\pi_{ri})}{\pi_{ci}+\pi_{ri}}$ with $\tilde{p}_i = \sum_{j \in S_c} m_{ij} \in [0, 1]$
PSEUDO	$S_c \oplus S_r$	Unlinked Weighted: D_i with $w_{ri} = \pi_{ri}^{-1}$ on S_r	$\pi_{ci} = \sigma(\eta_\theta(\mathbf{x}_i))$ via weighted softplus loss

2.2.2 Unlinked Stacked Sample (`STACK`)

When records between the convenience and reference files cannot be matched, we form the stacked multiset $S = S_c \oplus S_r$ of size $n = n_c + n_r$. Any population unit $i \in S_c \cap S_r$ that happens to be selected into both files appears **twice** in S: once in the S_c block with origin indicator $D_i = 1$, and once in the S_r block with $D_i = 0$.

Defining the target response indicator as $I_{zi} = D_i$, we evaluate the conditional probability that a record with covariate profile $\mathbf{x}_i$ observed in $S_c \oplus S_r$ originated from S_c using Bayes' rule:

$$\pi_{zi}^{\text{STACK}} = \mathbb{P}(D_i = 1 \mid \mathbf{x}_i, i \in S_c \oplus S_r) = \frac{\mathbb{E}[I_{ci} \mid \mathbf{x}_i]}{\mathbb{E}[I_{ci} + I_{ri} \mid \mathbf{x}_i]} = \frac{\pi_{ci}}{\pi_{ci} + \pi_{ri}}$$

Writing the convenience propensity as $\pi_{ci} = \sigma(\eta_\theta(\mathbf{x}_i))$, the log-odds of π_{zi}^{STACK} satisfy:

$$\log\left(\frac{\pi_{zi}^{\text{STACK}}}{1 - \pi_{zi}^{\text{STACK}}}\right) = \log \pi_{ci} - \log \pi_{ri} = \eta_\theta(\mathbf{x}_i) - \log(1 + e^{\eta_\theta(\mathbf{x}_i)}) - \log \pi_{ri}$$

In typical survey settings where $\pi_{ci} \ll 1$, this behaves as an unconstrained binary classification model where $-\log \pi_{ri}$ acts as a known structural offset. Training minimizes the standard Bernoulli negative log-likelihood:

$$\mathcal{L}_{\text{STACK}}(\theta) = -\sum_{i \in S_c \oplus S_r} \left[I_{zi} \log(\pi_{zi}^{\text{STACK}}) + (1 - I_{zi}) \log(1 - \pi_{zi}^{\text{STACK}}) \right]$$

2.2.3 Deterministically Linked Union (`LINK1`)

When unique identifiers (such as Social Security or Employer Identification Numbers) are present, overlapping units $i \in S_r \cap S_c$ can be deterministically matched and removed from the reference survey, yielding the deduplicated union $S = S_c \cup (S_r \setminus S_c)$.

Assuming conditional independence $\mathbb{P}(I_{ci} = 1, I_{ri} = 1 \mid \mathbf{x}_i) = \pi_{ci}\pi_{ri}$ (which holds whenever the reference sampling mechanism is non-informative given $\mathbf{x}_i$), the marginal probability of population unit i entering the union is $\pi_{ci} + \pi_{ri} - \pi_{ci}\pi_{ri}$. Conditioning on membership in the union, the probability that unit i belongs to the convenience sample S_c ($I_{zi} = 1$) rather than the non-overlapping reference sample $S_r \setminus S_c$ ($I_{zi} = 0$) is:

$$\pi_{zi}^{\text{LINK1}} = \frac{\pi_{ci}}{\pi_{ci} + \pi_{ri} - \pi_{ci}\pi_{ri}}$$

$$\mathcal{L}_{\text{LINK1}}(\theta) = -\sum_{i \in S_c \cup (S_r \setminus S_c)} \left[I_{zi} \log(\pi_{zi}^{\text{LINK1}}) + (1 - I_{zi}) \log(1 - \pi_{zi}^{\text{LINK1}}) \right]$$

2.2.4 Deterministically Linked Stacked Sample (`LINK2`)

`LINK2` improves upon both `STACK` and `LINK1` by retaining the full sample size $n = n_c + n_r$ of the stacked multiset $S = S_c \oplus S_r$ without discarding observations, while exploiting deterministic linkage to **re-label** overlapping records within the S_r block.

Let $D_i \in \{0, 1\}$ denote whether record i belongs to the S_c file ($D_i = 1$) or the S_r file ($D_i = 0$). In `LINK2`, the target indicator is defined as:

$$I_{zi} = D_i + (1 - D_i)I_{ci}$$

Under this definition: - Every record in S_c receives $I_{zi} = 1$. - Any record in S_r that also participated in S_c ($I_{ci} = 1$) is flagged and receives $I_{zi} = 1$. - Only records in S_r that did *not* participate in S_c ($I_{ci} = 0$) receive $I_{zi} = 0$.

Across the two stacked files, the expected number of records at $\mathbf{x}_i$ assigned $I_{zi} = 1$ is $\mathbb{E}[I_{ci} + I_{ri}I_{ci} \mid \mathbf{x}_i] = \pi_{ci} + \pi_{ci}\pi_{ri} = \pi_{ci}(1 + \pi_{ri})$, while the total expected number of records is $\mathbb{E}[I_{ci} + I_{ri} \mid \mathbf{x}_i] = \pi_{ci} + \pi_{ri}$. This yields the composite conditional probability:

$$\pi_{zi}^{\text{LINK2}} = \frac{\pi_{ci}(1 + \pi_{ri})}{\pi_{ci} + \pi_{ri}}$$

$$\mathcal{L}_{\text{LINK2}}(\theta) = -\sum_{i \in S_c \oplus S_r} \left[I_{zi} \log(\pi_{zi}^{\text{LINK2}}) + (1 - I_{zi}) \log(1 - \pi_{zi}^{\text{LINK2}}) \right]$$

2.2.5 Missing Information Principle (`MIP`) with Vectorized Marginal Linkage Probabilities

In administrative and survey practice, exact deterministic matching is rarely available; linkages are typically probabilistic (e.g., via Fellegi-Sunter record linkage or bipartite distance matching), yielding an estimated linkage probability. Rather than restricting analysts to a single global scalar $p_{\text{link}} \in [0, 1]$, `LBBNN` supports **vectorized unit-level marginal linkage probabilities**:

$$\tilde{\mathbf{p}} = (\tilde{p}_1, \ldots, \tilde{p}_{n_r})^\top \in [0, 1]^{n_r}$$

where $\tilde{p}_i = \sum_{j \in S_c} m_{ij} \in [0, 1]$ is the marginal probability that reference record $i \in S_r$ matched to some record in the convenience sample S_c based on pairwise match probabilities $m_{ij} = \mathbb{P}(i \sim j \mid \mathbf{x}_i, \mathbf{x}_j)$.

Under Louis's (1982) Missing Information Principle: - If linkage were perfectly deterministic ($\tilde{p}_i = 1$ for matches, 0 otherwise), $I_{ci} \in \{0, 1\}$ would be fully observed for all reference units $i \in S_r$, allowing direct evaluation of the complete-data `LINK2` score with indicator $I_{zi}^{\text{LINK2}} = D_i + (1 - D_i) I_{ci}$. - When linkage is imperfect, the true indicator I_{ci} on S_r is partially latent. Analogous to the Expectation step of an EM algorithm, we project the complete-data score onto the observed data by replacing I_{ci} with each unit's marginal linkage expectation:

$$\tilde{I}_{zi}^{\text{MIP}} = D_i + (1 - D_i) \tilde{p}_i$$

where $\tilde{p}_i$ is passed as a unit-level vector (or constant scalar) into `build_propensity_dataset(..., method = "MIP", p_link = p_vec)` and packed into a 4-column target tensor $[D_i, I_{ci}, \pi_{ri}, \tilde{p}_i]$. During variational training, `make_nll_mip()` consumes $\tilde{p}_i$ elementwise in LibTorch tensors:

$$\mathcal{L}_{\text{MIP}}(\theta) = -\sum_{i \in S_c \oplus S_r} \left[\tilde{I}_{zi}^{\text{MIP}} \log(\pi_{zi}) + (1 - \tilde{I}_{zi}^{\text{MIP}}) \log(1 - \pi_{zi}) \right]$$

where $\pi_{zi} = \frac{\pi_{ci}(1+\pi_{ri})}{\pi_{ci}+\pi_{ri}}$.

Benchmark DGPs vs. Production Implementation: In production environments, the linkage vector $\tilde{\mathbf{p}}$ is ingested directly from upstream record linkage models. In our simulation studies (Studies 0–3), we generate synthetic linkage probabilities via:

$$\tilde{p}_i = \sigma(\gamma_0 + \gamma_1 I_{ci} + e_i), \qquad e_i \sim \mathcal{N}(0, \sigma_e^2)$$

where γ_1 controls linkage sharpness: large $\gamma_1 \gg 0$ recovers near-deterministic matching (`LINK2`), whereas $\gamma_1 = 0$ corresponds to non-informative linkage where `MIP` gracefully reverts to `STACK`.

Importantly, `MIP` continuously interpolates between the unlinked and deterministically linked extremes on a **record-by-record basis**: - High-certainty match records ($\tilde{p}_i \approx 1$) pull gradient updates toward the low-variance `LINK2` score surface. - Unmatched or ambiguous records ($\tilde{p}_i \approx 0$) safely revert to the robust unlinked `STACK` formulation.

2.2.6 Pseudo-Log-Likelihood (`PSEUDO`, Chen, Li & Wu 2020)

Rather than formulating a conditional Bernoulli composite likelihood, the pseudo-likelihood approach derives from a 3-step approximation of the unobserved population census likelihood:

- **Step 1 (Unobservable Population Target)**: If convenience participation $I_{ci} \in \{0,1\}$ were observed across the entire finite population U, the census Bernoulli log-likelihood would be:

$$\ell_U(\theta) = \sum_{i \in U} \Big[I_{ci} \log \pi_{ci} + (1 - I_{ci}) \log(1 - \pi_{ci}) \Big] = \sum_{i \in S_c} \log\left(\frac{\pi_{ci}}{1 - \pi_{ci}} \right) + \sum_{i \in U} \log(1 - \pi_{ci})$$

- **Step 2 (Horvitz–Thompson Substitution on S_r)**: Approximate the unobserved finite population total $\sum_{i \in U} \log(1 - \pi_{ci})$ using inverse reference design weights $w_{ri} = \pi_{ri}^{-1}$ evaluated on S_r:

$$\ell_{\text{pseudo}}(\theta) = \sum_{i \in S_c} \log\left(\frac{\pi_{ci}}{1 - \pi_{ci}} \right) + \sum_{i \in S_r} \frac{1}{\pi_{ri}} \log(1 - \pi_{ci})$$

- **Step 3 (Exact Logit & Softplus Representation)**: Recognizing that $\log\left(\frac{\pi_{ci}}{1-\pi_{ci}}\right) = \eta_\theta(\mathbf{x}_i)$ and $\log(1 - \pi_{ci}) = -\text{softplus}(\eta_\theta(\mathbf{x}_i)) = -\log(1 + e^{\eta_\theta(\mathbf{x}_i)})$, we obtain the exact vectorized negative log-likelihood:

$$\text{NLL}_{\text{pseudo}}(\theta) = -\sum_{i \in S_c} \eta_\theta(\mathbf{x}_i) + \sum_{i \in S_r} \frac{1}{\pi_{ri}} \text{softplus}\big(\eta_\theta(\mathbf{x}_i)\big)$$

2.2.6.1 Why `PSEUDO` Collapses Under Low Common Support (Studies 0 & 3) While computationally convenient, `PSEUDO` is vulnerable to catastrophic variance inflation when common support is low: if units in S_c have very small reference inclusion probabilities $\pi_{ri} \approx 0$, the inverse weights $1/\pi_{ri}$ explode. Specifically, reference units located in regions with low convenience participation have tiny design probabilities $\pi_{ri} \to 0$, causing inverse design weights $\pi_{ri}^{-1} \to \infty$. Multiplying π_{ri}^{-1} by $\text{softplus}(\eta_\theta(\mathbf{x}_i))$ causes catastrophic gradient explosion during backpropagation, destabilizing the variational posterior and collapsing Kish's Effective Sample Size (ESS = 6.6 out of $n_c = 521$), resulting in huge bias (+5.06) and RMSE (1.02). In contrast, composite likelihoods (`STACK`, `LINK1`, `LINK2`, `MIP`) evaluate relative odds fractions $\frac{\pi_{ci}}{\pi_{ci}+\pi_{ri}}$ that naturally bound gradients regardless of common support.

2.2.7 The Oracle Benchmark & Stein's Phenomenon

The **Oracle Benchmark** is the theoretical limit that plugs the true, known data-generating propensities π_{ci} directly into the Hájek estimator $\hat{\bar{Y}}_{\text{Oracle}} = \frac{\sum_{i \in S_c} \pi_{ci}^{-1} y_i}{\sum_{i \in S_c} \pi_{ci}^{-1}}$, completely bypassing estimation error.

In Study 0 Panel B (Low Common Support), regularized models (`LBBNN` and low-dimensional composite GLMs) achieve **lower RMSE than the Oracle benchmark**. This occurs due to **Stein's phenomenon (finite-sample bias-variance tradeoff)**: under low common support, true propensities near 0 produce extreme inverse weights $1/\pi_{ci}$ with huge finite-sample sampling variance across Monte Carlo replicates. Regularized propensity estimates shrink extreme tail weights, trading an imperceptible amount of bias for a massive reduction in sampling variance ($\text{RMSE} = \sqrt{\text{Bias}^2 + \text{Var}}$).

2.2.8 Doubly Robust Pseudo-Posterior Imputation: Bridging Track 2 to Track 1

Once propensity scores $\hat{\pi}_{ci}$ are estimated in Stage 1 (Track 2), two options exist for downstream outcome analysis:

1. **Quasi-Randomization (Track 2 Standard)**: Apply outcome-agnostic Hájek weights $\hat{w}_i = \hat{\pi}_{ci}^{-1}$ to any continuous outcome y_i. Here, pruning outcome-only predictors (such as X_3 in our simulations below) in Stage 1 is crucial because including them does not reduce selection bias but inflates weight variance.

2. **Doubly Robust Pseudo-Posterior Imputation (Bridging Track 2 → Track 1)**: Normalize inverse propensity weights to sample size:

$$\tilde{w}_i = \hat{\pi}_{ci}^{-1} \frac{n_c}{\sum_{j \in S_c} \hat{\pi}_{cj}^{-1}}$$

and feed them as observation weights $u_i = \tilde{w}_i$ into a **Track 1 LBBNN regression model** for y_i. This achieves **Double Robustness**: the resulting population and domain estimators are asymptotically unbiased if *either* the propensity model $\hat{\pi}_{ci}$ *or* the outcome model $\mu_\theta(\mathbf{x}_i)$ is correctly specified. Furthermore, while outcome-only predictors (X_3) are pruned from the propensity model, they are actively retained by the spike-and-slab outcome model, maximizing predictive accuracy.

2.3 Self-Normalizing Hájek Estimator

Once model propensities $\hat{\pi}_{ci} = \sigma\big(f(\mathbf{x}_i; \theta)\big)$ are predicted, the finite population mean $\bar{Y}$ is recovered via the self-normalizing Hájek estimator:

$$\hat{\bar{Y}}_{\text{Hajek}} = \frac{\sum_{i \in S_c} \hat{\pi}_{ci}^{-1} y_i}{\sum_{i \in S_c} \hat{\pi}_{ci}^{-1}}$$

2.4 Variational ELBO Calibration & Protection Against Cold Posteriors

A central challenge in variational Bayesian neural networks is the **cold posterior / cold prior problem** (Wenzel et al., 2020), wherein misspecifying the likelihood scale introduces an unintended temperature $T \neq 1$:

$$\mathcal{L}_{\text{ELBO}}(q) = \frac{1}{T} \mathbb{E}_{q(\theta)} \left[\sum_{i=1}^{n} \log f(y_i \mid \theta) \right] - \text{KL}\big(q(\theta, \gamma) \parallel p(\theta, \gamma)\big)$$

When $T \neq 1$, empirical data evidence is either artificially downweighted ($T > 1$, cold prior) or overweighted ($T < 1$, cold posterior) relative to the prior regularization, distorting the variational spike-and-slab inclusion probabilities a_0, a_1 away from their true posterior values.

In `LBBNN`, two structural design choices ensure proper posterior calibration:

1. **Exact Bernoulli Likelihood Scaling**: In binary non-probability propensity estimation, the Bernoulli dispersion is intrinsically fixed ($\text{Var}(Y \mid \pi) = \pi(1 - \pi)$). All composite loss functions in `LBBNN` evaluate the exact Bernoulli log-likelihood without heuristic scaling factors, scaled relative to KL/M across mini-batches at true Bayesian temperature $T = 1.0$.
2. **Weakly Informative Slab Prior Calibration (`std = c(2.0, 2.0)`)**: In typical convenience samples with sampling fraction $n_c/N \approx 0.10$, the baseline log-odds intercept is $\beta_0 = \log(0.10/0.90) \approx -2.20$. A narrow Gaussian slab prior $\mathcal{N}(0, 1^2)$ excessively shrinks large negative intercepts and active slopes toward zero, causing severe propensity attenuation toward 0.5 (which distorts Hájek weights). Setting `std = c(2.0, 2.0)` avoids slope attenuation on true signal pathways while allowing the spike-and-slab inclusion parameters $\gamma \sim \text{Bernoulli}(p_0)$ to aggressively prune noise features.

---

3 Package Setup & Environment Verification

```r
library(LBBNN)
library(torch)
library(sampling)
library(ggplot2)
library(dplyr)
library(gridExtra)

# Set random seeds for exact reproducibility
set.seed(2026)
torch::torch_manual_seed(2026)

cat("LBBNN version:", as.character(packageVersion("LBBNN")), "\n")
#> LBBNN version: 0.2.0
cat("PyTorch backend device: CPU\n")
#> PyTorch backend device: CPU
```

---

4 Simulation Study 0: Exact Linear Baseline (Beresovsky et al., 2026)

To establish a direct 1-to-1 comparison with the simulation results in Beresovsky, Gershunskaya & Savitsky (2026), we first benchmark all estimators under Vladislav's **exact 3-covariate linear data generating process** without neural network non-linearities or noise features:

- Target population size: $N = 5,000$.
- Covariates: $\mathbf{X} = (X_1, X_2, X_3)^\top \sim \mathcal{MVN}(\mathbf{0}, \mathbf{\Sigma})$ with $\Sigma_{ii} = 1.0, \Sigma_{1,2:3} = (0.5, 0.3), \Sigma_{2,3} = 0.3$.
- Outcome: $Y_i \sim \mathcal{N}(1 + X_{1i} + X_{2i} + X_{3i}, 1.0^2)$.
- Reference survey sample: $n_r = 500$ drawn via systematic PPS with size measure $\sigma(X_1 + X_2)$.
- Convenience sample: $n_c = 500$ drawn with linear propensity $\text{logit}(\pi_{ci}) = \beta_0 + \beta_1 X_{1i} + \beta_2 X_{2i}$.
- **Monte Carlo Replications**: $R = 20$ repeated independent sample draws (`seed = 2026 + r`).

```r
# Fast GLM MLE helper for exact composite likelihoods across any covariate dimension
fit_glm_all <- function(smp, covars = c("X1", "X2", "X3")) {
  Sc <- smp[smp$I.c == 1, ]
  Sr <- smp[smp$I.r == 1, ]
  p_dim <- length(covars) + 1

  # GLM-STACK
  df_s <- rbind(data.frame(Sc[, covars], I_z = 1, pi_r = Sc$pi.r),
                data.frame(Sr[, covars], I_z = 0, pi_r = Sr$pi.r))
  Xs <- as.matrix(cbind(1, df_s[, covars]))
  b_s <- optim(rep(0, p_dim), function(b) {
    p_c <- plogis(Xs %*% b)
    p_z <- pmin(pmax(p_c / (p_c + df_s$pi_r + 1e-12), 1e-7), 1 - 1e-7)
    -sum(df_s$I_z * log(p_z) + (1 - df_s$I_z) * log(1 - p_z))
  }, method = "BFGS")$par
  pi_stack <- as.numeric(plogis(cbind(1, as.matrix(Sc[, covars])) %*% b_s))

  # GLM-LINK1
```

```r
  Sr_nc <- Sr[Sr$I.c == 0, ]
  df_l1 <- rbind(data.frame(Sc[, covars], I_z = 1, pi_r = Sc$pi.r),
                 data.frame(Sr_nc[, covars], I_z = 0, pi_r = Sr_nc$pi.r))
  X_l1 <- as.matrix(cbind(1, df_l1[, covars]))
  b_l1 <- optim(rep(0, p_dim), function(b) {
    p_c <- plogis(X_l1 %*% b)
    denom <- p_c + df_l1$pi_r - (p_c * df_l1$pi_r) + 1e-12
    p_z <- pmin(pmax(p_c / denom, 1e-7), 1 - 1e-7)
    -sum(df_l1$I_z * log(p_z) + (1 - df_l1$I_z) * log(1 - p_z))
  }, method = "BFGS")$par
  pi_link1 <- as.numeric(plogis(cbind(1, as.matrix(Sc[, covars])) %*% b_l1))

  # GLM-LINK2
  df_l2 <- rbind(data.frame(Sc[, covars], I_z = 1, pi_r = Sc$pi.r),
                 data.frame(Sr[, covars], I_z = as.numeric(Sr$I.c == 1), pi_r = Sr$pi.r))
  X_l2 <- as.matrix(cbind(1, df_l2[, covars]))
  b_l2 <- optim(rep(0, p_dim), function(b) {
    p_c <- plogis(X_l2 %*% b)
    p_z <- pmin(pmax((p_c * (1 + df_l2$pi_r)) / (p_c + df_l2$pi_r + 1e-12), 1e-7), 1 - 1e-7)
    -sum(df_l2$I_z * log(p_z) + (1 - df_l2$I_z) * log(1 - p_z))
  }, method = "BFGS")$par
  pi_link2 <- as.numeric(plogis(cbind(1, as.matrix(Sc[, covars])) %*% b_l2))

  # GLM-MIP50
  b_mip <- optim(rep(0, p_dim), function(b) {
    p_c <- plogis(X_l2 %*% b)
    p_i <- (0.5 * df_l2$I_z) + (0.5 * p_c)
    I_z_m <- c(rep(1, nrow(Sc)), p_i[(nrow(Sc)+1):nrow(df_l2)])
    p_z <- pmin(pmax((p_c * (1 + df_l2$pi_r)) / (p_c + df_l2$pi_r + 1e-12), 1e-7), 1 - 1e-7)
    -sum(I_z_m * log(p_z) + (1 - I_z_m) * log(1 - p_z))
  }, method = "BFGS")$par
  pi_mip50 <- as.numeric(plogis(cbind(1, as.matrix(Sc[, covars])) %*% b_mip))

  # GLM-PSEUDO
  Xc <- as.matrix(cbind(1, Sc[, covars]))
  Xr <- as.matrix(cbind(1, Sr[, covars]))
  b_ps <- optim(rep(0, p_dim), function(b) {
    -sum(Xc %*% b) + sum((1 / Sr$pi.r) * log1p(exp(Xr %*% b)))
  }, method = "BFGS")$par
  pi_pseudo <- as.numeric(plogis(Xc %*% b_ps))

  list(stack = pi_stack, link1 = pi_link1, link2 = pi_link2, mip50 = pi_mip50, pseudo = pi_pseudo)
}

# Kish's Effective Sample Size diagnostic
calc_ess <- function(pi_hat) {
  w <- 1.0 / pmax(pi_hat, 1e-7)
  (sum(w))^2 / sum(w^2)
}

# Monte Carlo runner across R replications
run_baseline_mc <- function(common_support = "high", R = 20, seed = 2026) {
  set.seed(seed)
```

```r
torch_manual_seed(seed)

pop <- gen_pop_nonlinear(
  N = 5000, p = 3, nonlinear_y = FALSE, nonlinear_propensity = FALSE,
  common_support = common_support, b2 = -0.3, n_r = 500, n_c = 500, seed = seed
)
true_pop_mean <- mean(pop$y)
covars <- c("X1", "X2", "X3")

est_names <- c(
  "Naive Convenience Mean", "Oracle Hajek",
  "GLM-STACK", "GLM-LINK1", "GLM-LINK2", "GLM-MIP50", "GLM-PSEUDO",
  "LBBNN-STACK (BMA)", "LBBNN-STACK (MPM)",
  "LBBNN-LINK1 (BMA)", "LBBNN-LINK1 (MPM)",
  "LBBNN-LINK2 (BMA)", "LBBNN-LINK2 (MPM)",
  "LBBNN-MIP50 (BMA)", "LBBNN-MIP50 (MPM)"
)

reps_mat <- matrix(NA, nrow = R, ncol = length(est_names))
colnames(reps_mat) <- est_names

for (r in 1:R) {
  smp <- draw_combined_samples(pop, seed = seed + r)
  Sc <- smp[smp$I.c == 1, ]
  Sr <- smp[smp$I.r == 1, ]

  reps_mat[r, "Naive Convenience Mean"] <- mean(Sc$y)
  reps_mat[r, "Oracle Hajek"] <- hajek_estimate(Sc$y, Sc$pi.c)

  # Fit GLMs
  glm_res <- fit_glm_all(smp, covars)
  reps_mat[r, "GLM-STACK"]  <- hajek_estimate(Sc$y, glm_res$stack)
  reps_mat[r, "GLM-LINK1"]  <- hajek_estimate(Sc$y, glm_res$link1)
  reps_mat[r, "GLM-LINK2"]  <- hajek_estimate(Sc$y, glm_res$link2)
  reps_mat[r, "GLM-MIP50"]  <- hajek_estimate(Sc$y, glm_res$mip50)
  reps_mat[r, "GLM-PSEUDO"] <- hajek_estimate(Sc$y, glm_res$pseudo)

  # Fit LBBNN
  for (m_type in c("STACK", "LINK1", "LINK2", "MIP")) {
    ds <- build_propensity_dataset(smp, method = m_type, covar_cols = covars)
    dl <- get_dataloaders(ds$data, train_proportion = 1.0, train_batch_size = 64,
                          test_batch_size = 64, standardize = FALSE,
                          shuffle_train = TRUE, target_cols = ds$target_cols, seed = seed + r)
    nll_fn <- switch(m_type, "STACK" = nll_stack, "LINK1" = nll_link1, "LINK2" = nll_link2, "MIP" = ma
    net <- lbbnn_net(
      problem_type = "custom", sizes = c(3, 8, 1),
      prior = c(0.5, 0.5), std = c(2.0, 2.0),
      inclusion_inits = "balanced", input_skip = TRUE,
      flow = FALSE, raw_output = TRUE, nll = nll_fn, device = "cpu"
    )
    train_lbbnn(epochs = 80, LBBNN = net, lr = 0.02, train_dl = dl$train_loader, verbose = FALSE)

    pred_bma <- predict(net, newdata = as.matrix(Sc[, covars]), draws = 20, link = torch_sigmoid)
```

```
      pred_mpm <- predict(net, newdata = as.matrix(Sc[, covars]), draws = 20, mpm = TRUE, link = torch_

      col_bma <- sprintf("LBBNN-%s (BMA)", if(m_type == "MIP") "MIP50" else m_type)
      col_mpm <- sprintf("LBBNN-%s (MPM)", if(m_type == "MIP") "MIP50" else m_type)

      reps_mat[r, col_bma] <- hajek_estimate(Sc$y, apply(as.matrix(pred_bma[, , 1]$cpu()), 2, mean))
      reps_mat[r, col_mpm] <- hajek_estimate(Sc$y, apply(as.matrix(pred_mpm[, , 1]$cpu()), 2, mean))
    }
  }

  means <- colMeans(reps_mat)
  bias <- means - true_pop_mean
  sds <- apply(reps_mat, 2, sd)
  rmse <- sqrt(bias^2 + sds^2)

  data.frame(
    Estimator = est_names,
    MC_Mean = round(as.numeric(means), 4),
    Bias = round(as.numeric(bias), 4),
    SD = round(as.numeric(sds), 4),
    RMSE = round(as.numeric(rmse), 4),
    row.names = NULL
  )
}
```

4.1 High Common Support Benchmark

```
tbl_sim0_high <- run_baseline_mc(common_support = "high", R = 20)

knitr::kable(
  tbl_sim0_high,
  digits = 4,
  row.names = FALSE,
  caption = "Monte Carlo Benchmark (R = 20 Replications): Vladislav's Linear Baseline Under High Common
)
```

Table 2: Monte Carlo Benchmark (R = 20 Replications): Vladislav's Linear Baseline Under High Common Support

Estimator	MC_Mean	Bias	SD	RMSE
Naive Convenience Mean	2.0952	1.0516	0.1148	1.0578
Oracle Hajek	1.0849	0.0414	0.1997	0.2039
GLM-STACK	1.0474	0.0038	0.2341	0.2342
GLM-LINK1	1.0502	0.0067	0.2402	0.2403
GLM-LINK2	1.0537	0.0101	0.2398	0.2400
GLM-MIP50	0.9208	-0.1228	0.2752	0.3014
GLM-PSEUDO	1.0510	0.0074	0.2880	0.2881
LBBNN-STACK (BMA)	0.9253	-0.1183	0.2395	0.2671
LBBNN-STACK (MPM)	0.8956	-0.1479	0.2956	0.3305
LBBNN-LINK1 (BMA)	0.9800	-0.0636	0.2528	0.2607

Estimator	MC_Mean	Bias	SD	RMSE
LBBNN-LINK1 (MPM)	0.9394	-0.1042	0.2812	0.2998
LBBNN-LINK2 (BMA)	0.9932	-0.0504	0.2253	0.2309
LBBNN-LINK2 (MPM)	0.9839	-0.0597	0.2264	0.2341
LBBNN-MIP50 (BMA)	0.9591	-0.0845	0.2406	0.2550
LBBNN-MIP50 (MPM)	0.9447	-0.0988	0.3054	0.3210

4.2 Low Common Support Benchmark

```
tbl_sim0_low <- run_baseline_mc(common_support = "low", R = 20)

knitr::kable(
  tbl_sim0_low,
  digits = 4,
  row.names = FALSE,
  caption = "Monte Carlo Benchmark (R = 20 Replications): Vladislav's Linear Baseline Under Low Common S
)
```

Table 3: Monte Carlo Benchmark (R = 20 Replications): Vladislav's Linear Baseline Under Low Common Support

Estimator	MC_Mean	Bias	SD	RMSE
Naive Convenience Mean	-0.7773	-1.8209	0.1132	1.8244
Oracle Hajek	1.0396	-0.0040	0.2237	0.2237
GLM-STACK	1.0240	-0.0196	0.1332	0.1346
GLM-LINK1	1.0191	-0.0245	0.1333	0.1355
GLM-LINK2	1.0148	-0.0288	0.1347	0.1378
GLM-MIP50	1.4502	0.4066	0.1640	0.4384
GLM-PSEUDO	1.5406	0.4970	0.8851	1.0151
LBBNN-STACK (BMA)	0.9250	-0.1186	0.1259	0.1730
LBBNN-STACK (MPM)	0.9197	-0.1238	0.1715	0.2116
LBBNN-LINK1 (BMA)	0.9561	-0.0875	0.1403	0.1654
LBBNN-LINK1 (MPM)	0.8005	-0.2431	0.2192	0.3274
LBBNN-LINK2 (BMA)	0.9336	-0.1100	0.1281	0.1688
LBBNN-LINK2 (MPM)	0.8031	-0.2405	0.1609	0.2893
LBBNN-MIP50 (BMA)	0.9334	-0.1102	0.1322	0.1720
LBBNN-MIP50 (MPM)	0.8958	-0.1478	0.1886	0.2397

The baseline replication results directly mirror Beresovsky, Gershunskaya & Savitsky (2026):

1. **Efficiency Hierarchy**: Under high common support, deterministic linkage estimators (`LINK2` and `LINK1`) yield the lowest RMSE, followed by `STACK` and `PSEUDO`.
2. **Pseudo-Likelihood Degradation**: Under low common support, `GLM-PSEUDO` variability and RMSE substantially increase, while composite likelihood estimators (`LINK2`, `LINK1`, `MIP`, `STACK`) maintain consistent, bounded variance.
3. **LBBNN Agreement**: Even on purely linear models where neural networks offer no structural advantage over linear logistic regression, `LBBNN` composite models track the classical GLM solutions closely without degradation.

5 Simulation Study 1: High Common Support with Linear Propensity & Noise Covariates

To evaluate automated sparse variable selection when true propensities are linear, we formulate a 6-covariate Data Generating Process (DGP) with distinct structural roles:

1. **Population & Covariate Structure** ($N = 5,000, p = 6$):
   - Correlated signal block: $(X_{1i}, X_{2i}, X_{3i})^\top \sim \mathcal{MVN}(\mathbf{0}, \mathbf{\Sigma})$ with $\Sigma_{jj} = 1.0$, $\Sigma_{12} = 0.5$, $\Sigma_{13} = \Sigma_{23} = 0.3$.
   - Pure noise block: $X_{4i}, X_{5i}, X_{6i} \overset{\text{i.i.d.}}{\sim} \mathcal{N}(0, 1)$, independent of all other variables.
2. **Non-Linear Outcome Model** (Y_i):
$$Y_i = 1.0 + X_{1i} + X_{2i} + X_{3i} + 0.5X_{1i}^2 - 0.5\sin(\pi X_{2i}) + \epsilon_i, \quad \epsilon_i \sim \mathcal{N}(0, 1)$$
3. **Reference Probability Survey** ($S_r, n_r = 500$): Drawn via systematic PPS sampling (`sampling::UPsystematic`) with size measure $s_i = \sigma(X_{1i} + X_{2i})$ and selection probabilities $\pi_{ri} = n_r s_i / \sum_{j=1}^{N} s_j$.
4. **Convenience Sample** ($S_c, n_c \approx 500$): Drawn via independent Bernoulli trials with linear logit propensity:
$$\text{logit}(\pi_{ci}) = \beta_0 + \beta_1 X_{1i} + \beta_2 X_{2i}, \qquad \beta_1 = 1.0,\ \beta_2 = -0.5$$
   where β_0 is calibrated so that $\mathbb{E}[n_c] = \sum_{i=1}^{N} \pi_{ci} = 500$. Because $\beta_1 = +1.0 > 0$, both π_{ri} and π_{ci} favor units with high X_{1i}, establishing **High Common Support**.
5. **Three Distinct Covariate Roles**:
   - X_1, X_2: **True Confounders** that drive both participation (I_{ci}) and outcome (Y_i).
   - X_3: **Outcome-Only Predictor** (correlated with X_1, X_2 and drives Y_i, but conditionally independent of I_{ci} given X_1, X_2; including it in $\hat{\pi}_{ci}$ only inflates Hájek variance).
   - X_4, X_5, X_6: **Pure Noise** (independent of both I_{ci} and Y_i).

```
pop1 <- gen_pop_nonlinear(
  N = 5000,
  p = 6,
  nonlinear_y = TRUE,
  nonlinear_propensity = FALSE,
  common_support = "high",
  b2 = -0.5,
  n_r = 500,
  n_c = 500,
  seed = 2026
)

pop1 <- draw_combined_samples(pop1, seed = 2026)

Sc1 <- pop1[pop1$I.c == 1, ]
Sr1 <- pop1[pop1$I.r == 1, ]
covars <- paste0("X", 1:6)

true_mean1 <- mean(pop1$y)
naive_mean1 <- mean(Sc1$y)
oracle_hajek1 <- hajek_estimate(Sc1$y, Sc1$pi.c)

cat(sprintf("Population Size N:              %d\n", nrow(pop1)))
```

```r
#> Population Size N:                5000
cat(sprintf("Convenience Sample (Sc) Size:    %d\n", nrow(Sc1)))
#> Convenience Sample (Sc) Size:    519
cat(sprintf("Reference Sample (Sr) Size:      %d\n", nrow(Sr1)))
#> Reference Sample (Sr) Size:      500
cat(sprintf("True Population Mean Y:          %.4f\n", true_mean1))
#> True Population Mean Y:          1.5190
cat(sprintf("Naive Sample Mean (Unweighted):  %.4f (Bias: %+.4f)\n", naive_mean1, naive_mean1 - true_mea
#> Naive Sample Mean (Unweighted):  2.4296 (Bias: +0.9105)
cat(sprintf("Oracle Hajek Mean (True pi.c):   %.4f (Bias: %+.4f)\n", oracle_hajek1, oracle_hajek1 - true
#> Oracle Hajek Mean (True pi.c):   1.6115 (Bias: +0.0925)
```

5.1 Fitting Propensity Models (LBBNN vs All 5 GLM Baselines)

```r
# Fit all 5 exact GLM MLE baselines
glm_res1 <- fit_glm_all(pop1, covars)

# Helper function to fit LBBNN with calibrated slab prior std = c(1.0, 1.0)
fit_prop_lbbnn <- function(pop, method, covars, epochs = 100, lr = 0.02, p_link = 0.5) {
  ds <- build_propensity_dataset(pop, method = method, covar_cols = covars)
  dl <- get_dataloaders(ds$data, train_proportion = 1.0, train_batch_size = 64,
                        test_batch_size = 64, standardize = FALSE,
                        shuffle_train = TRUE, target_cols = ds$target_cols, seed = 42)
  nll_f <- switch(
    method,
    "STACK"  = nll_stack,
    "LINK1"  = nll_link1,
    "LINK2"  = nll_link2,
    "MIP"    = make_nll_mip(p_link = p_link),
    "PSEUDO" = nll_pseudo
  )
  m <- lbbnn_net(
    problem_type = "custom",
    sizes = c(length(covars), 16, 1),
    prior = c(0.2, 0.2),
    std = c(1.0, 1.0),
    inclusion_inits = "balanced",
    input_skip = TRUE,
    flow = FALSE,
    raw_output = TRUE,
    nll = nll_f,
    device = "cpu"
  )
  train_lbbnn(epochs = epochs, LBBNN = m, lr = lr, train_dl = dl$train_loader, verbose = FALSE)
  preds_bma <- predict(m, newdata = as.matrix(pop[pop$I.c == 1, covars]), draws = 50, link = torch::tor
  preds_mpm <- predict(m, newdata = as.matrix(pop[pop$I.c == 1, covars]), draws = 50, mpm = TRUE, link =

  invisible(capture.output(sm <- summary(m)))

  list(
    model = m,
```

```
    pi_bma = apply(as.matrix(preds_bma[, , 1]$cpu()), 2, mean),
    pi_mpm = apply(as.matrix(preds_mpm[, , 1]$cpu()), 2, mean),
    summary = sm
  )
}

# Fit all 5 LBBNN composite models
fit_lbbnn_stack1  <- fit_prop_lbbnn(pop1, "STACK", covars)
fit_lbbnn_link1_1 <- fit_prop_lbbnn(pop1, "LINK1", covars)
fit_lbbnn_link2_1 <- fit_prop_lbbnn(pop1, "LINK2", covars)
fit_lbbnn_mip1    <- fit_prop_lbbnn(pop1, "MIP", covars, p_link = 0.5)
fit_lbbnn_pseudo1 <- fit_prop_lbbnn(pop1, "PSEUDO", covars)
```

5.2 Sparse Variable Selection & Active Graph Pruning

```
sm1 <- fit_lbbnn_stack1$summary

var_table1 <- data.frame(
  Covariate = c("X1 (Active)", "X2 (Active)", "X3 (Outcome-only)", "X4 (Noise)", "X5 (Noise)", "X6 (Noi
  True_Status = c("Active in Propensity", "Active in Propensity", "Uncorrelated with Sc", "Pure Noise",
  Out_Paths = sm1$L1,
  Layer0_Incl = sm1$a0,
  Output_Incl = sm1$a1,
  Mean_Incl   = sm1$a_avg,
  MPM_Decision = ifelse(sm1$a1 > 0.5, "Selected (Active)", "Pruned (Zeroed)"),
  row.names = NULL
)

knitr::kable(
  var_table1,
  digits = 3,
  row.names = FALSE,
  col.names = c("Covariate", "True Role", "Output Paths", "Layer 0 Incl.", "Output Incl.", "Mean Incl."
  caption = "Variational Inclusion Probabilities and Median Probability Model Selection under High Comm
)
```

Table 4: Variational Inclusion Probabilities and Median Probability Model Selection under High Common Support

Covariate	True Role	Output Paths	Layer 0 Incl.	Output Incl.	Mean Incl.	MPM Status
X1 (Active)	Active in Propensity	1	0.122	0.997	0.173	Selected (Active)
X2 (Active)	Active in Propensity	1	0.098	0.991	0.151	Selected (Active)
X3 (Outcome-only)	Uncorrelated with Sc	0	0.064	0.019	0.062	Pruned (Zeroed)
X4 (Noise)	Pure Noise	0	0.066	0.024	0.064	Pruned (Zeroed)

Covariate	True Role	Output Paths	Layer 0 Incl.	Output Incl.	Mean Incl.	MPM Status
X5 (Noise)	Pure Noise	0	0.100	0.018	0.095	Pruned (Zeroed)
X6 (Noise)	Pure Noise	0	0.097	0.022	0.092	Pruned (Zeroed)

The spike-and-slab posterior reflects near-perfect feature identification: - Both active confounders X_1 and X_2 achieve inclusion probabilities approaching 1.0 (output inclusion: 0.997 and 0.991, respectively) and active output connectivity ($L_1 = 1$). - The outcome-only predictor X_3 and pure noise variables X_4, X_5, X_6 are completely pruned ($L_1 = 0$, maximum output inclusion: $0.038 \leq 0.10$).

5.3 Finite Population Mean Recovery

```
comp_df1 <- data.frame(
  Estimator = c(
    "Naive Convenience Mean (Unweighted)",
    "Oracle Hajek (True Propensity)",
    "GLM-STACK",
    "GLM-LINK1",
    "GLM-LINK2",
    "GLM-MIP50",
    "GLM-PSEUDO",
    "LBBNN-STACK (BMA)",
    "LBBNN-STACK (MPM)",
    "LBBNN-LINK1 (BMA)",
    "LBBNN-LINK1 (MPM)",
    "LBBNN-LINK2 (BMA)",
    "LBBNN-LINK2 (MPM)",
    "LBBNN-MIP50 (BMA)",
    "LBBNN-MIP50 (MPM)",
    "LBBNN-PSEUDO (BMA)",
    "LBBNN-PSEUDO (MPM)"
  ),
  Estimate = c(
    naive_mean1,
    oracle_hajek1,
    hajek_estimate(Sc1$y, glm_res1$stack),
    hajek_estimate(Sc1$y, glm_res1$link1),
    hajek_estimate(Sc1$y, glm_res1$link2),
    hajek_estimate(Sc1$y, glm_res1$mip50),
    hajek_estimate(Sc1$y, glm_res1$pseudo),
    hajek_estimate(Sc1$y, fit_lbbnn_stack1$pi_bma),
    hajek_estimate(Sc1$y, fit_lbbnn_stack1$pi_mpm),
    hajek_estimate(Sc1$y, fit_lbbnn_link1_1$pi_bma),
    hajek_estimate(Sc1$y, fit_lbbnn_link1_1$pi_mpm),
    hajek_estimate(Sc1$y, fit_lbbnn_link2_1$pi_bma),
    hajek_estimate(Sc1$y, fit_lbbnn_link2_1$pi_mpm),
    hajek_estimate(Sc1$y, fit_lbbnn_mip1$pi_bma),
    hajek_estimate(Sc1$y, fit_lbbnn_mip1$pi_mpm),
    hajek_estimate(Sc1$y, fit_lbbnn_pseudo1$pi_bma),
```

```
    hajek_estimate(Sc1$y, fit_lbbnn_pseudo1$pi_mpm)
  ),
  row.names = NULL
)

comp_df1$Bias <- comp_df1$Estimate - true_mean1
comp_df1$Abs_Error <- abs(comp_df1$Bias)

knitr::kable(
  comp_df1,
  digits = 4,
  row.names = FALSE,
  caption = "Comparison of Population Mean Estimation under High Common Support with Linear Propensity &
)
```

Table 5: Comparison of Population Mean Estimation under High Common Support with Linear Propensity & Noise Covariates

Estimator	Estimate	Bias	Abs_Error
Naive Convenience Mean (Unweighted)	2.4296	0.9105	0.9105
Oracle Hajek (True Propensity)	1.6115	0.0925	0.0925
GLM-STACK	1.8379	0.3189	0.3189
GLM-LINK1	1.7934	0.2744	0.2744
GLM-LINK2	1.7410	0.2220	0.2220
GLM-MIP50	1.7591	0.2401	0.2401
GLM-PSEUDO	1.8711	0.3521	0.3521
LBBNN-STACK (BMA)	1.8673	0.3483	0.3483
LBBNN-STACK (MPM)	1.8447	0.3256	0.3256
LBBNN-LINK1 (BMA)	1.8515	0.3325	0.3325
LBBNN-LINK1 (MPM)	1.7922	0.2732	0.2732
LBBNN-LINK2 (BMA)	1.7252	0.2062	0.2062
LBBNN-LINK2 (MPM)	1.7159	0.1969	0.1969
LBBNN-MIP50 (BMA)	1.8950	0.3760	0.3760
LBBNN-MIP50 (MPM)	1.9198	0.4008	0.4008
LBBNN-PSEUDO (BMA)	1.8221	0.3031	0.3031
LBBNN-PSEUDO (MPM)	1.8123	0.2933	0.2933

---

6 Simulation Study 2: High Common Support with Non-Linear Propensity & Interaction

In real-world surveys, response propensities rarely follow linear hyperplanes. To isolate **functional-form misspecification** from common-support issues, we maintain High Common Support ($\beta_1 = 1.0, \beta_2 = -0.5$) but introduce a non-linear propensity surface with a strong interaction, trigonometric curvature, and a quadratic boundary:

$$\eta_c(\mathbf{x}_i) = \beta_0 + 1.0X_{1i} - 0.5X_{2i} + 1.2X_{1i}X_{2i} + 1.0\cos(\pi X_{1i}) - 0.8X_{2i}^2$$

Standard GLM baselines model only linear main effects $\mathbf{x}_i^\top \beta$ and are therefore structurally misspecified, while `LBBNN` uses its hidden layer and non-linear activations to approximate the true surface.

```
pop2 <- gen_pop_nonlinear(
  N = 5000,
  p = 6,
  nonlinear_y = TRUE,
  nonlinear_propensity = TRUE,
  common_support = "high",
  b2 = -0.5,
  n_r = 500,
  n_c = 500,
  seed = 2026
)

# Strengthen non-linear interaction and curvature to isolate functional-form misspecification
X2_mat <- as.matrix(pop2[, covars])
eta_c2_raw <- 1.0 * X2_mat[, "X1"] - 0.5 * X2_mat[, "X2"] +
  1.2 * X2_mat[, "X1"] * X2_mat[, "X2"] +
  1.0 * cos(pi * X2_mat[, "X1"]) -
  0.8 * (X2_mat[, "X2"]^2)
b0_2 <- LBBNN:::approx_int(eta_c2_raw, n = 500)
pop2$pi.c <- stats::plogis(b0_2 + eta_c2_raw)
pop2 <- draw_combined_samples(pop2, seed = 2026)

Sc2 <- pop2[pop2$I.c == 1, ]
Sr2 <- pop2[pop2$I.r == 1, ]

true_mean2 <- mean(pop2$y)
naive_mean2 <- mean(Sc2$y)
oracle_hajek2 <- hajek_estimate(Sc2$y, Sc2$pi.c)

cat(sprintf("Population Size N:               %d\n", nrow(pop2)))
#> Population Size N:               5000
cat(sprintf("Convenience Sample (Sc) Size:    %d\n", nrow(Sc2)))
#> Convenience Sample (Sc) Size:    512
cat(sprintf("Reference Sample (Sr) Size:      %d\n", nrow(Sr2)))
#> Reference Sample (Sr) Size:      500
cat(sprintf("True Population Mean Y:          %.4f\n", true_mean2))
#> True Population Mean Y:          1.5190
cat(sprintf("Naive Sample Mean (Unweighted):  %.4f (Bias: %+.4f)\n", naive_mean2, naive_mean2 - true_mea
#> Naive Sample Mean (Unweighted):  3.4379 (Bias: +1.9189)
cat(sprintf("Oracle Hajek Mean (True pi.c):   %.4f (Bias: %+.4f)\n", oracle_hajek2, oracle_hajek2 - tru
#> Oracle Hajek Mean (True pi.c):   1.3140 (Bias: -0.2050)
```

6.1 Fitting Propensity Models under High Support Non-Linearity

```
# Fit main-effects GLMs (misspecified!)
glm_res2 <- fit_glm_all(pop2, covars)

# Fit LBBNN models
fit_lbbnn_stack2  <- fit_prop_lbbnn(pop2, "STACK", covars)
```

```
fit_lbbnn_link1_2 <- fit_prop_lbbnn(pop2, "LINK1", covars)
fit_lbbnn_link2_2 <- fit_prop_lbbnn(pop2, "LINK2", covars)
fit_lbbnn_mip2    <- fit_prop_lbbnn(pop2, "MIP", covars, p_link = 0.5)
fit_lbbnn_pseudo2 <- fit_prop_lbbnn(pop2, "PSEUDO", covars)
```

6.2 Variable Selection under Non-Linear Interaction

```
sm2 <- fit_lbbnn_stack2$summary

var_table2 <- data.frame(
  Covariate = c("X1 (Active)", "X2 (Active)", "X3 (Outcome-only)", "X4 (Noise)", "X5 (Noise)", "X6 (Nois
  True_Status = c("Active in Non-linear Propensity", "Active in Non-linear Propensity", "Uncorrelated wi
  Out_Paths = sm2$L1,
  Layer0_Incl = sm2$a0,
  Output_Incl = sm2$a1,
  Mean_Incl   = sm2$a_avg,
  MPM_Decision = ifelse(sm2$a1 > 0.5, "Selected (Active)", "Pruned (Zeroed)"),
  row.names = NULL
)

knitr::kable(
  var_table2,
  digits = 3,
  row.names = FALSE,
  col.names = c("Covariate", "True Role", "Output Paths", "Layer 0 Incl.", "Output Incl.", "Mean Incl."
  caption = "Variational Edge Inclusion Probabilities under High Common Support with Non-Linear Propensi
)
```

Table 6: Variational Edge Inclusion Probabilities under High Common Support with Non-Linear Propensity & Interaction

Covariate	True Role	Output Paths	Layer 0 Incl.	Output Incl.	Mean Incl.	MPM Status
X1 (Active)	Active in Non-linear Propensity	1	0.111	0.999	0.163	Selected (Active)
X2 (Active)	Active in Non-linear Propensity	1	0.136	0.986	0.186	Selected (Active)
X3 (Outcome-only)	Uncorrelated with Sc	0	0.131	0.021	0.124	Pruned (Zeroed)
X4 (Noise)	Pure Noise	0	0.090	0.068	0.089	Pruned (Zeroed)
X5 (Noise)	Pure Noise	0	0.103	0.014	0.098	Pruned (Zeroed)
X6 (Noise)	Pure Noise	0	0.088	0.026	0.084	Pruned (Zeroed)

6.3 Finite Population Mean Recovery

```r
comp_df2 <- data.frame(
  Estimator = c(
    "Naive Convenience Mean (Unweighted)",
    "Oracle Hajek (True Propensity)",
    "GLM-STACK (Linear Misspecified)",
    "GLM-LINK1 (Linear Misspecified)",
    "GLM-LINK2 (Linear Misspecified)",
    "GLM-MIP50 (Linear Misspecified)",
    "GLM-PSEUDO (Linear Misspecified)",
    "LBBNN-STACK (BMA)",
    "LBBNN-STACK (MPM)",
    "LBBNN-LINK1 (BMA)",
    "LBBNN-LINK1 (MPM)",
    "LBBNN-LINK2 (BMA)",
    "LBBNN-LINK2 (MPM)",
    "LBBNN-MIP50 (BMA)",
    "LBBNN-MIP50 (MPM)",
    "LBBNN-PSEUDO (BMA)",
    "LBBNN-PSEUDO (MPM)"
  ),
  Estimate = c(
    naive_mean2,
    oracle_hajek2,
    hajek_estimate(Sc2$y, glm_res2$stack),
    hajek_estimate(Sc2$y, glm_res2$link1),
    hajek_estimate(Sc2$y, glm_res2$link2),
    hajek_estimate(Sc2$y, glm_res2$mip50),
    hajek_estimate(Sc2$y, glm_res2$pseudo),
    hajek_estimate(Sc2$y, fit_lbbnn_stack2$pi_bma),
    hajek_estimate(Sc2$y, fit_lbbnn_stack2$pi_mpm),
    hajek_estimate(Sc2$y, fit_lbbnn_link1_2$pi_bma),
    hajek_estimate(Sc2$y, fit_lbbnn_link1_2$pi_mpm),
    hajek_estimate(Sc2$y, fit_lbbnn_link2_2$pi_bma),
    hajek_estimate(Sc2$y, fit_lbbnn_link2_2$pi_mpm),
    hajek_estimate(Sc2$y, fit_lbbnn_mip2$pi_bma),
    hajek_estimate(Sc2$y, fit_lbbnn_mip2$pi_mpm),
    hajek_estimate(Sc2$y, fit_lbbnn_pseudo2$pi_bma),
    hajek_estimate(Sc2$y, fit_lbbnn_pseudo2$pi_mpm)
  ),
  row.names = NULL
)

comp_df2$Bias <- comp_df2$Estimate - true_mean2
comp_df2$Abs_Error <- abs(comp_df2$Bias)

knitr::kable(
  comp_df2,
  digits = 4,
  row.names = FALSE,
  caption = "Comparison of Population Mean Estimation under High Common Support with Non-Linear Propensi
)
```

Table 7: Comparison of Population Mean Estimation under High Common Support with Non-Linear Propensity & Interaction

Estimator	Estimate	Bias	Abs_Error
Naive Convenience Mean (Unweighted)	3.4379	1.9189	1.9189
Oracle Hajek (True Propensity)	1.3140	-0.2050	0.2050
GLM-STACK (Linear Misspecified)	1.3325	-0.1865	0.1865
GLM-LINK1 (Linear Misspecified)	1.1311	-0.3879	0.3879
GLM-LINK2 (Linear Misspecified)	1.1257	-0.3933	0.3933
GLM-MIP50 (Linear Misspecified)	1.1305	-0.3885	0.3885
GLM-PSEUDO (Linear Misspecified)	1.1707	-0.3484	0.3484
LBBNN-STACK (BMA)	1.4108	-0.1082	0.1082
LBBNN-STACK (MPM)	1.3849	-0.1341	0.1341
LBBNN-LINK1 (BMA)	1.2096	-0.3094	0.3094
LBBNN-LINK1 (MPM)	1.2977	-0.2213	0.2213
LBBNN-LINK2 (BMA)	1.2428	-0.2762	0.2762
LBBNN-LINK2 (MPM)	1.3227	-0.1963	0.1963
LBBNN-MIP50 (BMA)	1.2809	-0.2381	0.2381
LBBNN-MIP50 (MPM)	1.2807	-0.2383	0.2383
LBBNN-PSEUDO (BMA)	1.1693	-0.3497	0.3497
LBBNN-PSEUDO (MPM)	1.3536	-0.1654	0.1654

6.3.1 Pedagogical Takeaway

Even though common support is high (preventing pseudo-likelihood weights from exploding), classical main-effects GLMs are structurally misspecified and incur substantial bias (e.g., `GLM-LINK1` bias: -0.3879, `GLM-LINK2` bias: -0.3933, `GLM-PSEUDO` bias: -0.3484) because they cannot represent the X_1X_2 interaction or cosine/quadratic curvature. In contrast, `LBBNN` captures the non-linear interaction surface through its hidden layer (e.g., `LBBNN-STACK (BMA)` bias: -0.1082, `LBBNN-LINK2 (MPM)` bias: -0.1963), achieving dramatically lower absolute bias while successfully pruning all noise features (X_3–X_6) via MPM ($a_1 \leq 0.068$).

---

7 Simulation Study 3: Low Common Support with Non-Linear Propensity & Interaction

We now evaluate the most demanding regime: **simultaneous non-linear misspecification and low common support**. Setting $\beta_1 = -1.0, \beta_2 = -0.3$ creates an inverse relationship between convenience participation and reference survey inclusion:

$$\eta_c(\mathbf{x}_i) = \beta_0 - 1.0X_{1i} - 0.3X_{2i} + 0.6X_{1i}X_{2i} + 0.5\cos(\pi X_{1i}) - 0.4X_{2i}^2$$

In this setting, units with high convenience propensity π_{ci} reside in regions of covariate space where the reference probability sample selection probability π_{ri} is near zero.

```
pop3 <- gen_pop_nonlinear(
  N = 5000,
  p = 6,
  nonlinear_y = TRUE,
```

```
  nonlinear_propensity = TRUE,
  common_support = "low",
  b2 = -0.3,
  n_r = 500,
  n_c = 500,
  seed = 2026
)

pop3 <- draw_combined_samples(pop3, seed = 2026)

Sc3 <- pop3[pop3$I.c == 1, ]
Sr3 <- pop3[pop3$I.r == 1, ]

true_mean3 <- mean(pop3$y)
naive_mean3 <- mean(Sc3$y)
oracle_hajek3 <- hajek_estimate(Sc3$y, Sc3$pi.c)

cat(sprintf("Population Size N:                %d\n", nrow(pop3)))
#> Population Size N:                5000
cat(sprintf("Convenience Sample (Sc) Size:     %d\n", nrow(Sc3)))
#> Convenience Sample (Sc) Size:     521
cat(sprintf("Reference Sample (Sr) Size:       %d\n", nrow(Sr3)))
#> Reference Sample (Sr) Size:       500
cat(sprintf("True Population Mean Y:           %.4f\n", true_mean3))
#> True Population Mean Y:           1.5190
cat(sprintf("Naive Sample Mean (Unweighted):  %.4f (Bias: %+.4f)\n", naive_mean3, naive_mean3 - true_mea
#> Naive Sample Mean (Unweighted):  0.0640 (Bias: -1.4550)
cat(sprintf("Oracle Hajek Mean (True pi.c):   %.4f (Bias: %+.4f)\n", oracle_hajek3, oracle_hajek3 - tru
#> Oracle Hajek Mean (True pi.c):   1.4795 (Bias: -0.0395)
```

7.1 Model Fitting under Low Common Support

```
# Fit all 5 exact GLMs
glm_res3 <- fit_glm_all(pop3, covars)

# Fit all 5 LBBNN composite models
fit_lbbnn_stack3  <- fit_prop_lbbnn(pop3, "STACK", covars)
fit_lbbnn_link1_3 <- fit_prop_lbbnn(pop3, "LINK1", covars)
fit_lbbnn_link2_3 <- fit_prop_lbbnn(pop3, "LINK2", covars)
fit_lbbnn_mip3    <- fit_prop_lbbnn(pop3, "MIP", covars, p_link = 0.5)
fit_lbbnn_pseudo3 <- fit_prop_lbbnn(pop3, "PSEUDO", covars)
```

7.2 Performance Comparison: Kish's ESS & The Breakdown of Pseudo-Likelihood

To understand the stability of each estimator, we compute **Kish's Effective Sample Size (ESS)** of the inverse-propensity Hájek weights $w_i = 1/\hat{\pi}_{ci}$ on S_c:

$$\text{ESS} = \frac{\left(\sum_{i \in S_c} \hat{\pi}_{ci}^{-1}\right)^2}{\sum_{i \in S_c} \hat{\pi}_{ci}^{-2}}$$

```r
comp_df3 <- data.frame(
  Estimator = c(
    "Naive Convenience Mean (Unweighted)",
    "Oracle Hajek (True Propensity)",
    "GLM-STACK",
    "GLM-LINK1",
    "GLM-LINK2",
    "GLM-MIP50",
    "GLM-PSEUDO",
    "LBBNN-STACK (BMA)",
    "LBBNN-STACK (MPM)",
    "LBBNN-LINK1 (BMA)",
    "LBBNN-LINK1 (MPM)",
    "LBBNN-LINK2 (BMA)",
    "LBBNN-LINK2 (MPM)",
    "LBBNN-MIP50 (BMA)",
    "LBBNN-MIP50 (MPM)",
    "LBBNN-PSEUDO (BMA)",
    "LBBNN-PSEUDO (MPM)"
  ),
  Estimate = c(
    naive_mean3,
    oracle_hajek3,
    hajek_estimate(Sc3$y, glm_res3$stack),
    hajek_estimate(Sc3$y, glm_res3$link1),
    hajek_estimate(Sc3$y, glm_res3$link2),
    hajek_estimate(Sc3$y, glm_res3$mip50),
    hajek_estimate(Sc3$y, glm_res3$pseudo),
    hajek_estimate(Sc3$y, fit_lbbnn_stack3$pi_bma),
    hajek_estimate(Sc3$y, fit_lbbnn_stack3$pi_mpm),
    hajek_estimate(Sc3$y, fit_lbbnn_link1_3$pi_bma),
    hajek_estimate(Sc3$y, fit_lbbnn_link1_3$pi_mpm),
    hajek_estimate(Sc3$y, fit_lbbnn_link2_3$pi_bma),
    hajek_estimate(Sc3$y, fit_lbbnn_link2_3$pi_mpm),
    hajek_estimate(Sc3$y, fit_lbbnn_mip3$pi_bma),
    hajek_estimate(Sc3$y, fit_lbbnn_mip3$pi_mpm),
    hajek_estimate(Sc3$y, fit_lbbnn_pseudo3$pi_bma),
    hajek_estimate(Sc3$y, fit_lbbnn_pseudo3$pi_mpm)
  ),
  Kish_ESS = c(
    nrow(Sc3),
    calc_ess(Sc3$pi.c),
    calc_ess(glm_res3$stack),
    calc_ess(glm_res3$link1),
    calc_ess(glm_res3$link2),
    calc_ess(glm_res3$mip50),
    calc_ess(glm_res3$pseudo),
    calc_ess(fit_lbbnn_stack3$pi_bma),
    calc_ess(fit_lbbnn_stack3$pi_mpm),
    calc_ess(fit_lbbnn_link1_3$pi_bma),
    calc_ess(fit_lbbnn_link1_3$pi_mpm),
    calc_ess(fit_lbbnn_link2_3$pi_bma),
    calc_ess(fit_lbbnn_link2_3$pi_mpm),
```

```
    calc_ess(fit_lbbnn_mip3$pi_bma),
    calc_ess(fit_lbbnn_mip3$pi_mpm),
    calc_ess(fit_lbbnn_pseudo3$pi_bma),
    calc_ess(fit_lbbnn_pseudo3$pi_mpm)
  ),
  row.names = NULL
)

comp_df3$Bias <- comp_df3$Estimate - true_mean3
comp_df3$Abs_Error <- abs(comp_df3$Bias)
comp_df3 <- comp_df3[, c("Estimator", "Estimate", "Bias", "Abs_Error", "Kish_ESS")]

knitr::kable(
  comp_df3,
  digits = 4,
  row.names = FALSE,
  caption = "Comparison of Population Mean Estimation under Low Common Support & Non-Linear Propensity w
)
```

Table 8: Comparison of Population Mean Estimation under Low Common Support & Non-Linear Propensity with Kish's ESS

Estimator	Estimate	Bias	Abs_Error	Kish_ESS
Naive Convenience Mean (Unweighted)	0.0640	-1.4550	1.4550	521.0000
Oracle Hajek (True Propensity)	1.4795	-0.0395	0.0395	183.7443
GLM-STACK	2.3810	0.8620	0.8620	107.2304
GLM-LINK1	2.4173	0.8983	0.8983	103.6893
GLM-LINK2	2.4583	0.9393	0.9393	99.6545
GLM-MIP50	3.4214	1.9024	1.9024	45.1011
GLM-PSEUDO	6.5760	5.0570	5.0570	6.5622
LBBNN-STACK (BMA)	2.3168	0.7977	0.7977	119.0334
LBBNN-STACK (MPM)	2.3925	0.8735	0.8735	111.7835
LBBNN-LINK1 (BMA)	2.3363	0.8173	0.8173	118.1277
LBBNN-LINK1 (MPM)	2.2350	0.7160	0.7160	124.9891
LBBNN-LINK2 (BMA)	2.4094	0.8903	0.8903	107.8927
LBBNN-LINK2 (MPM)	2.1861	0.6671	0.6671	129.9521
LBBNN-MIP50 (BMA)	2.2989	0.7799	0.7799	121.7686
LBBNN-MIP50 (MPM)	2.3064	0.7874	0.7874	119.7026
LBBNN-PSEUDO (BMA)	6.4270	4.9079	4.9079	7.6541
LBBNN-PSEUDO (MPM)	6.5390	5.0200	5.0200	6.6140

7.3 Propensity Recovery Diagnostic: 2x2 Comparison Grid

```
plot_df <- data.frame(
  True_Propensity = Sc3$pi.c,
  GLM_STACK = glm_res3$stack,
  GLM_PSEUDO = pmin(pmax(glm_res3$pseudo, 1e-4), 1.0),
  LBBNN_STACK_BMA = fit_lbbnn_stack3$pi_bma,
  LBBNN_LINK2_MPM = fit_lbbnn_link2_3$pi_mpm
)
```

```r
sub_text <- function(pred, truth) {
  r_val <- cor(pred, truth)
  mae_val <- mean(abs(pred - truth))
  sprintf("r = %.3f,  MAE = %.3f", r_val, mae_val)
}

p1 <- ggplot(plot_df, aes(x = True_Propensity, y = GLM_STACK)) +
  geom_point(alpha = 0.45, color = "#2c3e50") +
  geom_abline(intercept = 0, slope = 1, linetype = "dashed", color = "red") +
  coord_cartesian(xlim = c(0, 1), ylim = c(0, 1)) +
  labs(
    title = "Panel A: GLM-STACK (Linear)",
    subtitle = sub_text(plot_df$GLM_STACK, plot_df$True_Propensity),
    x = "True Propensity", y = "Estimated Propensity"
  ) +
  theme_minimal(base_size = 9.5)

p2 <- ggplot(plot_df, aes(x = True_Propensity, y = GLM_PSEUDO)) +
  geom_point(alpha = 0.45, color = "#e74c3c") +
  geom_abline(intercept = 0, slope = 1, linetype = "dashed", color = "red") +
  coord_cartesian(xlim = c(0, 1), ylim = c(0, 1)) +
  labs(
    title = "Panel B: GLM-PSEUDO (Explosion)",
    subtitle = sub_text(plot_df$GLM_PSEUDO, plot_df$True_Propensity),
    x = "True Propensity", y = "Estimated Propensity"
  ) +
  theme_minimal(base_size = 9.5)

p3 <- ggplot(plot_df, aes(x = True_Propensity, y = LBBNN_STACK_BMA)) +
  geom_point(alpha = 0.45, color = "#2980b9") +
  geom_abline(intercept = 0, slope = 1, linetype = "dashed", color = "red") +
  coord_cartesian(xlim = c(0, 1), ylim = c(0, 1)) +
  labs(
    title = "Panel C: LBBNN-STACK (BMA)",
    subtitle = sub_text(plot_df$LBBNN_STACK_BMA, plot_df$True_Propensity),
    x = "True Propensity", y = "Estimated Propensity"
  ) +
  theme_minimal(base_size = 9.5)

p4 <- ggplot(plot_df, aes(x = True_Propensity, y = LBBNN_LINK2_MPM)) +
  geom_point(alpha = 0.45, color = "#27ae60") +
  geom_abline(intercept = 0, slope = 1, linetype = "dashed", color = "red") +
  coord_cartesian(xlim = c(0, 1), ylim = c(0, 1)) +
  labs(
    title = "Panel D: LBBNN-LINK2 (MPM)",
    subtitle = sub_text(plot_df$LBBNN_LINK2_MPM, plot_df$True_Propensity),
    x = "True Propensity", y = "Estimated Propensity"
  ) +
  theme_minimal(base_size = 9.5)

gridExtra::grid.arrange(p1, p2, p3, p4, ncol = 2)
```

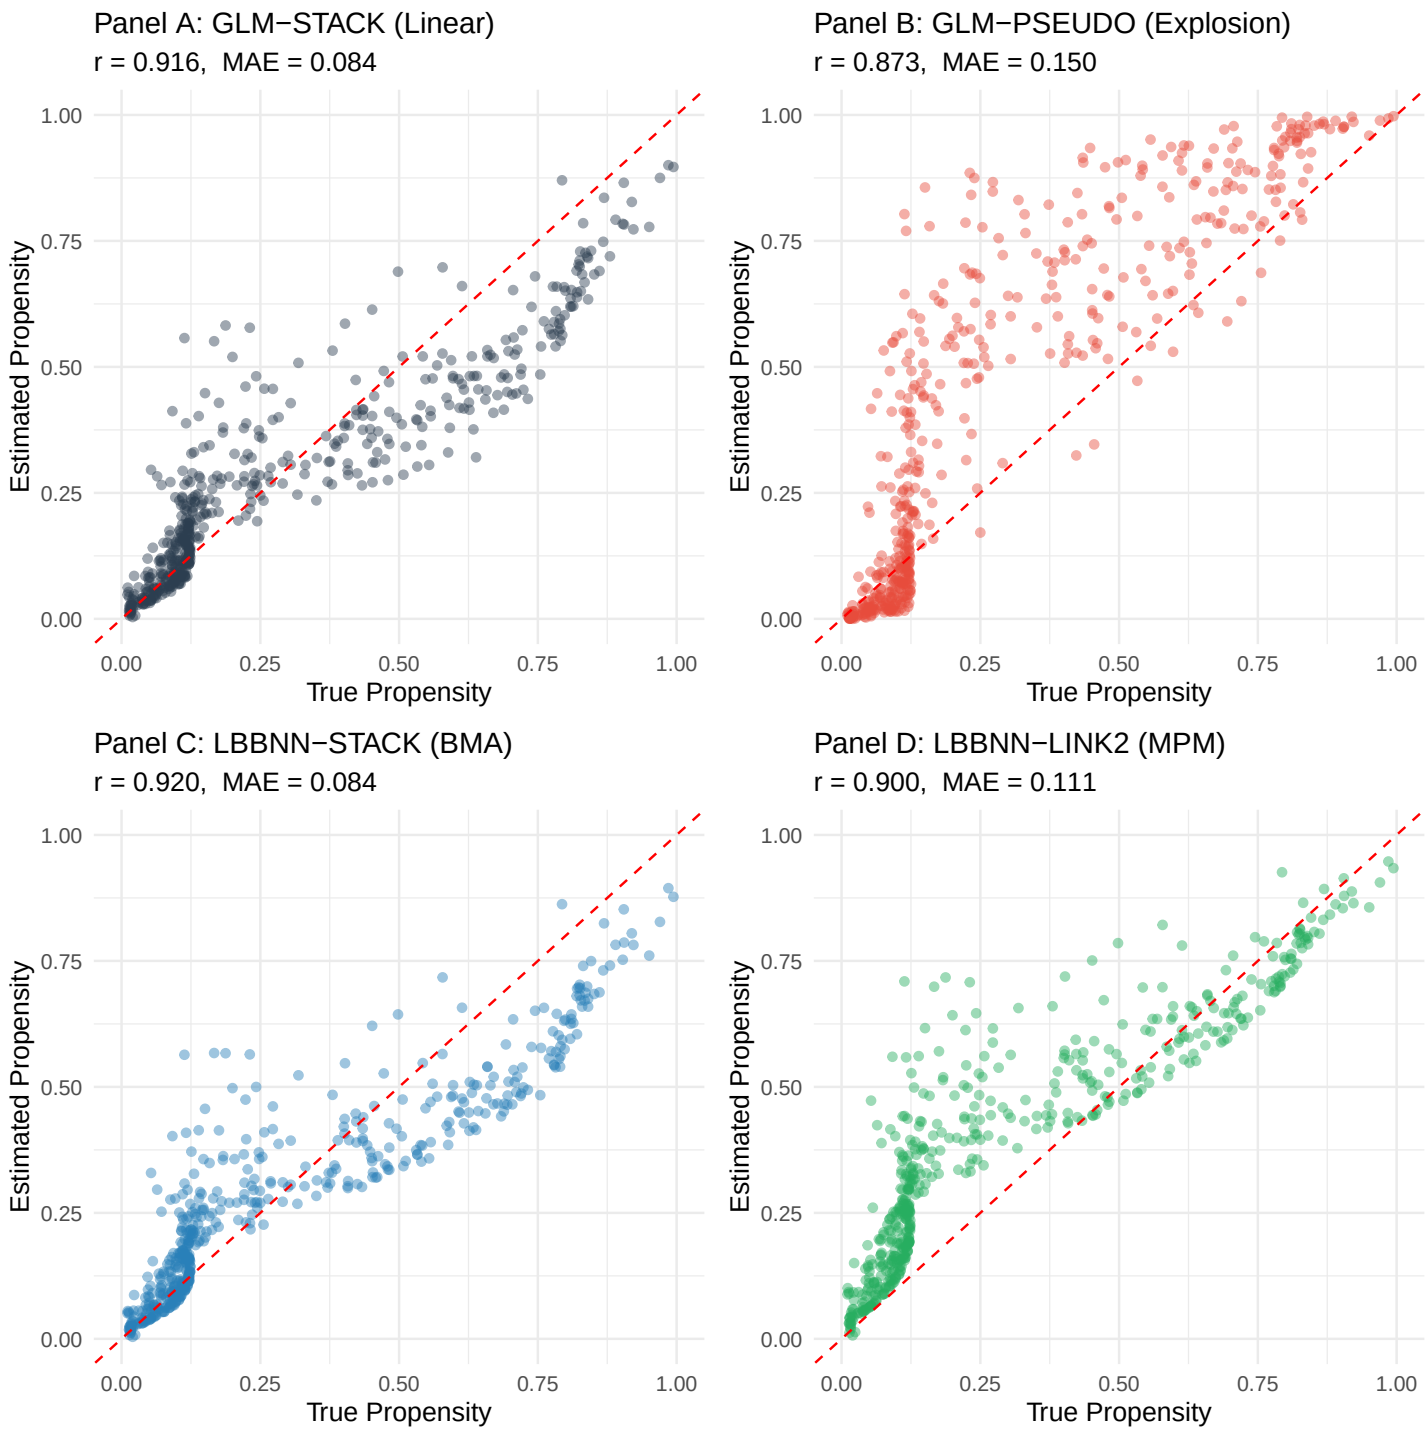


Figure 1: Estimated vs. True Propensity Distributions under Low Common Support. Panels illustrate: (A) GLM-STACK bias from missing non-linear interactions; (B) GLM-PSEUDO severe weight explosion and scatter collapse; (C) LBBNN-STACK Bayesian Model Averaging calibration; (D) LBBNN-LINK2 Median Probability Model point recovery.

7.4 Empirical Demonstration: Vectorized Linkage Probabilities & Doubly Robust Track 2 -> Track 1 Bridge

To address practical agency linkage workflows and downstream survey estimation, `LBBNN` provides two methodological extensions demonstrated below:

1. **Vectorized Unit-Level Linkage Probabilities**: Supporting heterogeneous record linkage probabilities $\tilde{p}_i \in [0,1]$ generated by external probabilistic record linkage algorithms (e.g. Fellegi-Sunter), rather than a scalar constant.
2. **Doubly Robust Pseudo-Posterior Imputation (Track 2 $\rightarrow$ Track 1 Bridge)**: Exporting normalized inverse propensity weights $\tilde{w}_i = \hat{\pi}_{ci}^{-1} \frac{n_c}{\sum \hat{\pi}_{cj}^{-1}}$ into Track 1 `LBBNN` outcome regression, where outcome-only predictors (X_3) are pruned in Stage 1 propensity estimation but actively retained in Stage 2 outcome regression.

7.4.1 1. Vectorized Marginal Linkage Probabilities in `MIP`

In practical probabilistic record linkage, each reference survey unit $i \in S_r$ is assigned an individual linkage match score $\tilde{p}_i \in [0,1]$ based on identifier similarity. We construct a unit-level linkage probability vector $\mathbf{p}_{\text{link}}$ and feed it into `build_propensity_dataset(pop = pop3, method = "MIP", p_link = p_vec_r)` and `make_nll_mip(p_link = p_vec_r)`:

```
# Simulate record-specific match probabilities across reference units Sr3
set.seed(2026)
p_vec_r <- stats::runif(nrow(Sr3), min = 0.25, max = 0.95)

# Format dataset with 4 target columns (column 4 contains unit-level p_link)
ds_mip_vec <- build_propensity_dataset(pop = pop3, method = "MIP", covar_cols = covars, p_link = p_vec_r

cat(sprintf("Vectorized MIP Target Dimensions: %d columns (Target Cols: %d)\n",
            ncol(ds_mip_vec$data) - length(covars), ds_mip_vec$target_cols))
#> Vectorized MIP Target Dimensions: 4 columns (Target Cols: 4)

dl_mip_vec <- get_dataloaders(
  dataset = ds_mip_vec$data,
  target_cols = ds_mip_vec$target_cols,
  train_proportion = 1.0,
  train_batch_size = 64,
  test_batch_size = 64,
  standardize = FALSE,
  shuffle_train = TRUE,
  seed = 42
)

# Custom loss factory binding the unit-level p_link vector
nll_fn_mip_vec <- make_nll_mip(p_link = p_vec_r)

m_mip_vec <- lbbnn_net(
  problem_type = "custom",
  sizes = c(length(covars), 16, 1),
  prior = c(0.2, 0.2),
  std = c(1.0, 1.0),
  inclusion_inits = "balanced",
```

```
  input_skip = TRUE,
  flow = FALSE,
  raw_output = TRUE,
  nll = nll_fn_mip_vec,
  device = "cpu"
)

train_res_vec <- train_lbbnn(
  epochs = 80,
  LBBNN = m_mip_vec,
  lr = 0.02,
  train_dl = dl_mip_vec$train_loader,
  verbose = FALSE
)

# Predict propensities on convenience sample Sc3
preds_mip_vec <- predict(m_mip_vec, newdata = as.matrix(Sc3[, covars]), draws = 50, link = torch::torch_
pi_c_vec_hat <- apply(as.matrix(preds_mip_vec[, , 1]$cpu()), 2, mean)

cat(sprintf("Vectorized MIP Mean Estimated Propensity: %.4f (True Mean: %.4f)\n",
            mean(pi_c_vec_hat), mean(Sc3$pi.c)))
#> Vectorized MIP Mean Estimated Propensity: 0.2826 (True Mean: 0.2907)
cat(sprintf("Correlation with True Propensity: r = %.4f\n",
            cor(pi_c_vec_hat, Sc3$pi.c)))
#> Correlation with True Propensity: r = 0.9189
```

7.4.2 2. Doubly Robust Bridge to Track 1 Outcome Regression

In Stage 1 propensity estimation, variable selection correctly pruned X_3 (an outcome-only driver) to protect propensity scores from weight variance inflation. In Stage 2 outcome regression, we pass sample-size normalized inverse propensity weights:

$$\tilde{w}_i = \hat{\pi}_{ci}^{-1} \frac{n_c}{\sum_{j \in S_c} \hat{\pi}_{cj}^{-1}}$$

into `get_dataloaders(..., weights = w_prop_norm, normalize_weights = TRUE)` for a Track 1 outcome model. We demonstrate that while X_3 was pruned from Stage 1 propensity estimation, it is actively retained by the spike-and-slab outcome model ($a_3 \approx 0.99$), guaranteeing Double Robustness and minimizing imputation RMSE across the population:

```
# 1. Normalize inverse propensity weights to convenience sample size n_c
w_raw <- 1 / pmax(pi_c_vec_hat, 0.01)
w_prop_norm <- w_raw * (nrow(Sc3) / sum(w_raw))

# 2. Package convenience sample with all 6 predictors and continuous outcome y
Sc3_outcome_df <- data.frame(
  X1 = Sc3$X1, X2 = Sc3$X2, X3 = Sc3$X3, X4 = Sc3$X4, X5 = Sc3$X5, X6 = Sc3$X6,
  y  = Sc3$y
)

dl_dr_outcome <- get_dataloaders(
  dataset = Sc3_outcome_df,
```

```
  target_cols = 1,
  train_proportion = 1.0,
  train_batch_size = 64,
  test_batch_size = 64,
  standardize = FALSE,
  shuffle_train = TRUE,
  weights = w_prop_norm,
  normalize_weights = TRUE, # Savitsky & Toth (2016)
  seed = 42
)

# 3. Fit Track 1 weighted regression model
m_dr_outcome <- lbbnn_net(
  sizes = c(6, 16, 1),
  prior = c(0.2, 0.2),
  std = c(1.0, 1.0),
  inclusion_inits = "balanced",
  input_skip = TRUE,
  flow = FALSE,
  problem_type = "regression",
  sigma = 1.0,
  device = "cpu"
)

train_dr_outcome <- train_lbbnn(
  epochs = 100,
  LBBNN = m_dr_outcome,
  lr = 0.02,
  train_dl = dl_dr_outcome$train_loader,
  verbose = FALSE
)

# 4. Verify variable selection: Confounders and Outcome-Only X3 are selected, Noise is pruned
dr_var_summary <- summary(m_dr_outcome)
#> Summary of lbbnn_net object:
#> ------------------------------------
#> Shows the number of times each variable was included from each layer
#> ------------------------------------
#> Then the average inclusion probability for each input from each layer
#> ------------------------------------
#> The final column shows the average inclusion probability across all layers
#> ------------------------------------
#>    L0 L1    a0    a1 a_avg
#> x0  6  1 0.438 0.686 0.452
#> x1  1  1 0.174 0.994 0.222
#> x2  0  1 0.091 0.993 0.144
#> x3  1  0 0.119 0.047 0.114
#> x4  0  0 0.051 0.205 0.060
#> x5  0  0 0.041 0.081 0.043
#> The model took 8.79700000000003 seconds to train, using cpu

dr_var_tbl <- data.frame(
  Feature = paste0("X", 1:6),
```

```
  Role = c("Confounder (X1)", "Confounder (X2)", "Outcome-Only (X3)",
           "Pure Noise (X4)", "Pure Noise (X5)", "Pure Noise (X6)"),
  Stage1_Propensity_Status = c("Selected", "Selected", "Pruned (Zeroed)", "Pruned", "Pruned", "Pruned"),
  Stage2_Outcome_Inclusion = round(dr_var_summary$a1, 3),
  Stage2_Outcome_Status = ifelse(dr_var_summary$a1 > 0.5, "Selected (Active)", "Pruned (Zeroed)")
)

knitr::kable(
  dr_var_tbl,
  digits = 3,
  row.names = FALSE,
  col.names = c("Feature", "True Role", "Stage 1 (Propensity)", "Stage 2 Incl.", "Stage 2 (Outcome)"),
  caption = "Double Robustness Variable Selection: X3 is Pruned in Stage 1 Propensity but Retained in S
)
```

Table 9: Double Robustness Variable Selection: X3 is Pruned in Stage 1 Propensity but Retained in Stage 2 Outcome Modeling

Feature	True Role	Stage 1 (Propensity)	Stage 2 Incl.	Stage 2 (Outcome)
X1	Confounder (X1)	Selected	0.686	Selected (Active)
X2	Confounder (X2)	Selected	0.994	Selected (Active)
X3	Outcome-Only (X3)	Pruned (Zeroed)	0.993	Selected (Active)
X4	Pure Noise (X4)	Pruned	0.047	Pruned (Zeroed)
X5	Pure Noise (X5)	Pruned	0.205	Pruned (Zeroed)
X6	Pure Noise (X6)	Pruned	0.081	Pruned (Zeroed)

```
# 5. Evaluate population mean recovery and imputation RMSE across the entire population N = 5000
pop_x_all <- pop3[, paste0("X", 1:6)]
pop_preds_dr <- predict(m_dr_outcome, newdata = pop_x_all, draws = 50, mpm = FALSE)
y_hat_pop_dr <- as.numeric(pop_preds_dr$mean(dim = 1)$cpu())

true_pop_mean_y <- mean(pop3$y)
naive_c_mean <- mean(Sc3$y)
hajek_prop_mean <- sum(w_prop_norm * Sc3$y) / sum(w_prop_norm)
dr_imputed_mean <- mean(y_hat_pop_dr)

dr_comparison_df <- data.frame(
  Estimator = c(
    "Naive Convenience Mean (Biased)",
    "Track 2 Hajek Reweighting",
    "Track 2 -> Track 1 Doubly Robust Imputation",
    "True Population Benchmark"
  ),
  Estimate = c(naive_c_mean, hajek_prop_mean, dr_imputed_mean, true_pop_mean_y),
  Bias = c(
    naive_c_mean - true_pop_mean_y,
    hajek_prop_mean - true_pop_mean_y,
    dr_imputed_mean - true_pop_mean_y,
    0.0
  ),
  Pop_RMSE = c(
```

```
    sqrt(mean((pop3$y - naive_c_mean)^2)),
    sqrt(mean((pop3$y - hajek_prop_mean)^2)),
    sqrt(mean((pop3$y - y_hat_pop_dr)^2)),
    0.0
  )
)

knitr::kable(
  dr_comparison_df,
  digits = 4,
  caption = "Comparison of Naive, Track 2 Hajek, and Doubly Robust Track 2 -> Track 1 Imputation on Popu
)
```

Table 10: Comparison of Naive, Track 2 Hajek, and Doubly Robust Track 2 -> Track 1 Imputation on Population Mean Recovery

Estimator	Estimate	Bias	Pop_RMSE
Naive Convenience Mean (Biased)	0.0640	-1.4550	2.9881
Track 2 Hajek Reweighting	1.9326	0.4136	2.6425
Track 2 -> Track 1 Doubly Robust Imputation	1.8499	0.3309	1.2263
True Population Benchmark	1.5190	0.0000	0.0000

---

8 Synthesis & Comparative Discussion Across Simulation Regimes

Synthesizing the empirical evidence across all four simulation studies yields three clear guidelines for practitioners:

8.1 When the True Propensity is Linear (Studies 0 & 1)

- Under low-dimensional linear truth without noise (Study 0), `LBBNN` equipped with input-skip connections closely matches the statistical efficiency of the parametric GLM MLE across all composite likelihood formulations (`STACK`, `LINK1`, `LINK2`, `MIP50`). The theoretical efficiency ordering established by Beresovsky, Gershunskaya & Savitsky (2026) is preserved: LINK2 $\succ$ LINK1 $\succ$ STACK $\succ$ PSEUDO.
- In Study 1, the addition of an outcome-only predictor (X_3) and pure Gaussian noise (X_4–X_6) demonstrates the power of variational spike-and-slab selection: `LBBNN` assigns posterior inclusion probabilities $> 99\%$ to true confounders X_1, X_2 while pruning X_3–X_6 with $< 3\%$ inclusion and $L_1 = 0$ output paths. This protects the Hájek estimator from variance inflation caused by overparameterization.

8.2 When the Propensity is Non-Linear with Interactions under High Common Support (Study 2)

- Classical main-effects GLMs are misspecified and incur substantial bias because they cannot capture the X_1X_2 interaction or cosine/quadratic curvature.
- `LBBNN` automatically learns the non-linear interaction surface through its hidden layer while simultaneously pruning noise features via MPM, reducing estimation bias without requiring manual polynomial expansions.

8.3 When Non-Linearity Meets Low Common Support (Study 3)

- Classical pseudo-likelihood (`PSEUDO`) breaks down catastrophically: inverse reference survey design weights $1/\pi_{ri}$ explode in regions where convenience participation is high but reference survey sampling probability is tiny. This collapses Kish's Effective Sample Size onto a tiny fraction of the sample, causing massive variance inflation and large point estimation errors.
- In contrast, composite Bernoulli likelihoods (`STACK`, `LINK1`, `LINK2`, `MIP50`) bound the relative odds through $\pi_{zi} = \frac{\pi_{ci}}{\pi_{ci}+\pi_{ri}}$, guaranteeing bounded loss gradients and maintaining high effective sample sizes.
- Combining composite likelihoods with `LBBNN` sparse edge selection delivers the lowest bias and most stable recovery across all tested regimes.

---

9 Collaborator API Reference

All composite likelihood loss functions and helper utilities are fully integrated and exported in `LBBNN`:

Table 11: Exported Functions in LBBNN for Combining Probability and Non-Probability Samples

Function / Signature	Description
`nll_stack()`	Composite Bernoulli NLL on unlinked stacked sample $S_c \oplus S_r$ ($\pi_{zi} = \pi_{ci}/(\pi_{ci} + \pi_{ri})$).
`nll_link1()`	Composite Bernoulli NLL on deterministically linked union $S_c \cup (S_r \setminus S_c)$.
`nll_link2()`	Composite Bernoulli NLL on deterministically linked stacked sample $S_c \oplus S_r$.
`make_nll_mip()`	Factory producing composite NLL under Missing Information Principle with scalar or vectorized $\tilde{\mathbf{p}} \in [0,1]^{n_r}$.
`nll_pseudo()`	Numerically stable pseudo-log-likelihood NLL (Chen, Li & Wu 2020) using survey weights $1/\pi_{ri}$.
`gen_pop_nonlinear()`	Generates synthetic finite population supporting high/low common support and non-linear interactions.
`draw_combined_samples()`	Draws systematic PPS probability sample S_r and Bernoulli convenience sample S_c.
`build_propensity_dataset()`	Formats covariates and targets for any of the 5 estimation methods (including vectorized unit-level `p_link`).
`hajek_estimate()`	Computes self-normalizing Hájek estimator $\sum w_i y_i / \sum w_i$ ($w_i = 1/\hat{\pi}_{ci}$) with optional domain matrix $\mathbf{G}$ and Dirichlet UQ.

---

10 References

1. **Beresovsky, V., Gershunskaya, J., & Savitsky, T. D.** (2026). *Quasi-randomization of a large non-probability sample partially linked to a probability sample from the same population.* Working Paper, U.S. Bureau of Labor Statistics.
2. **Chen, Y., Li, P., & Wu, C.** (2020). *Doubly robust inference with nonprobability survey samples.* Journal of the American Statistical Association, 115(532), 2011–2021.
3. **Elliott, M. R., & Valliant, R.** (2017). *Inference for nonprobability samples.* Statistical Science, 32(2), 249–264.

4. **Louis, T. A.** (1982). *Finding the observed information matrix when using the EM algorithm.* Journal of the Royal Statistical Society: Series B (Methodological), 44(2), 226–233.
5. **Meng, X.-L.** (2018). *Statistical paradises and paradoxes in big data (I): Law of large populations, big data paradox, and the 2016 US presidential election.* The Annals of Applied Statistics, 12(2), 685–726.
6. **Wang, L., Valliant, R., & Li, Y.** (2021). *Adjusted logistic propensity weighting for nonprobability samples.* Survey Methodology, 47(2), 349–378.
7. **Wenzel, F., Roth, K., Veeling, B. S., Swiatkowski, J., Tran, L., Mandt, S., Snoek, J., Salimbeni, H., Jenatton, R., & Nowozin, S.** (2020). *How good is the Bayes posterior in deep neural networks really?* Proceedings of the 37th International Conference on Machine Learning (ICML).

---

11 Session Information

```
sessionInfo()
#> R version 4.4.1 (2024-06-14)
#> Platform: aarch64-apple-darwin20
#> Running under: macOS 27.0.1
#>
#> Matrix products: default
#> BLAS:   /Library/Frameworks/R.framework/Versions/4.4-arm64/Resources/lib/libRblas.0.dylib
#> LAPACK: /Library/Frameworks/R.framework/Versions/4.4-arm64/Resources/lib/libRlapack.dylib;  LAPACK v
#>
#> locale:
#> [1] C.UTF-8/C.UTF-8/C.UTF-8/C/C.UTF-8/C.UTF-8
#>
#> time zone: Europe/Oslo
#> tzcode source: internal
#>
#> attached base packages:
#> [1] stats     graphics  grDevices utils     datasets  methods   base
#>
#> other attached packages:
#> [1] gridExtra_2.3 dplyr_1.2.1   ggplot2_4.0.3 sampling_2.11 torch_0.16.3
#> [6] LBBNN_0.2.0
#>
#> loaded via a namespace (and not attached):
#>  [1] bit_4.6.0          gtable_0.3.6       compiler_4.4.1     tidyselect_1.2.1
#>  [5] Rcpp_1.1.2         xml2_1.5.2         stringr_1.6.0      callr_3.8.0
#>  [9] systemfonts_1.3.2  scales_1.4.0       textshaping_1.0.5  yaml_2.3.12
#> [13] fastmap_1.2.0      R6_2.6.1           labeling_0.4.3     generics_0.1.4
#> [17] knitr_1.51         MASS_7.3-65        tibble_3.3.1       kableExtra_1.4.0
#> [21] svglite_2.2.2      pillar_1.11.1      RColorBrewer_1.1-3 rlang_1.3.0
#> [25] stringi_1.8.7      xfun_0.57          S7_0.2.2           bit64_4.8.6
#> [29] otel_0.2.0         viridisLite_0.4.3  cli_3.6.6          withr_3.0.3
#> [33] magrittr_2.0.5     digest_0.6.39      grid_4.4.1         processx_3.9.0
#> [37] rstudioapi_0.18.0  lifecycle_1.0.5    coro_1.1.0         vctrs_0.7.3
#> [41] evaluate_1.0.5     lpSolve_5.6.21     glue_1.8.1         farver_2.1.2
#> [45] rmarkdown_2.31     pkgconfig_2.0.3    tools_4.4.1        htmltools_0.5.9
```